

# Day 1: Overview of Core Concepts

---

ME314: Introduction to Data Science and Machine Learning

11 July 2022

LSE Summer School

Overview of Core Concepts

Supervised Learning

Unsupervised Learning

Machine Learning

# Overview of Core Concepts

---

## **Overview of Core Concepts**

- It is important to understand the ideas behind the various techniques, in order to know how and when to use them.
- One has to understand the simpler methods first, in order to grasp the more sophisticated ones.
- It is important to accurately assess the performance of a method, to know how well or how badly it is working (simpler methods often perform as well as fancier ones!).
- This is an exciting research area, having important applications in science, industry and policy.
- Machine learning is a fundamental ingredient in the training of a modern data scientist.

# Two main approaches to machine learning

- Supervised Learning
- Unsupervised Learning

# The Supervised Learning Problem

Starting point:

- Outcome measurement  $Y$  (also called dependent variable, response, target).
- Vector of  $p$  predictor measurements  $X$  (also called inputs, regressors, covariates, features, independent variables).
- In the **regression problem**,  $Y$  is quantitative (e.g price, blood pressure).
- In the **classification problem**,  $Y$  takes values in a finite, unordered set (survived/died, digit 0-9, cancer class of tissue sample).
- We have training data  $(x_1, y_1), \dots, (x_N, y_N)$ . These are observations (examples, instances) of these measurements.

# Objectives

On the basis of the training data we would like to:

- Accurately predict unseen test cases.
- Understand which inputs affect the outcome, and how.
- Assess the quality of our predictions and inferences.



# Unsupervised Learning

- No outcome variable, just a set of predictors (features) measured on a set of samples.
- Objective is more fuzzy – find groups of samples that behave similarly, find features that behave similarly, find linear combinations of features with the most variation.
- Difficult to know how well you are doing.
- Different from supervised learning, but can be useful as a pre-processing step for supervised learning.

# Machine Learning

---

## **Machine learning and this course**

# Machine learning

- **Machine learning** refers to a vast set of tools for *understanding data*.
- For a quantitative response  $Y$  and a set of predictors  $X$ :

$$Y = f(X) + \epsilon$$

- Here,  $f$  represents the *systematic* information that  $X$  provides about  $Y$ .
- Statistical learning refers to a set of approaches for estimating  $f$ .
- Most of the course we'll spend talking about different ways to estimate  $f$  and how to evaluate whether we've done a good job with it.

# Where does this course fit in?

- Supervised versus unsupervised learning.
- Regression versus classification.
- No single *best* method. We'll spend a lot of time choosing the most appropriate tool for a given dataset using different measures of the quality of fit. E.g. MSE

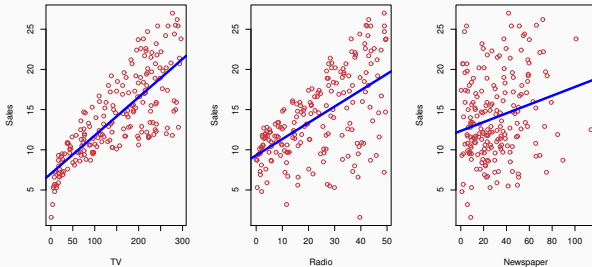
$$MSE_{training} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

$$MSE_{test} = \text{Ave}(\hat{f}(x_0) - y_0)^2$$

# Why should we bother with $f$ ?

1. **Prediction:**  $\hat{Y} = \hat{f}(X)$ , where  $\hat{f}$  is a *black box*.
2. **Inference:** How  $Y$  is changing as a function of  $X$ .
  - Depending on whether the ultimate goal is prediction, inference or a mix of both, we may deploy different methods for estimating  $f$ .
  - Also depending on the ultimate goal you may or may not care about evaluating the causal relationship between  $Y$  and  $X$ .

# What is Machine Learning?



- Shown are Sales vs TV, Radio and Newspaper, with a blue linear-regression line fit separately to each.
- We can predict Sales using a model

$$\text{Sales} \approx f(\text{TV}, \text{Radio}, \text{Newspaper})$$

# Notation

- Here Sales is a *response* or *target* that we wish to predict. We generically refer to the response as  $Y$ .
- TV is a *feature*, or *input*, or *predictor*; we name it  $X_1$ .
- Likewise name Radio as  $X_2$ , and so on.
- We can refer to the input vector collectively as

$$X = (X_1, X_2, X_3)$$

- Now we write our model as

$$Y = f(X) + \epsilon$$

where  $\epsilon$  captures measurement errors and other discrepancies.



## What is $f(X)$ good for?

- With a good  $f$  we can make predictions of  $Y$  at new points  $X = x$ .
- We can understand which components of  $X = (X_1, X_2, \dots, X_p)$  are important in explaining  $Y$ , and which are irrelevant. For example, Seniority and Years of Education have a big impact on Income, but Marital Status typically does not.
- Depending on the complexity of  $f$ , we may be able to understand how each component  $X_j$  of  $X$  affects  $Y$ .

- Is there an ideal  $f(X)$ ? In particular, what is a good value for  $f(X)$  at any selected value of  $X$ , say  $X = 4$ ?
- There can be many  $Y$  values at  $X = 4$ . A good value is

$$f(4) = E(Y|X = 4)$$

- $E(Y|X = 4)$  means expected value (average) of  $Y$  given  $X = 4$ .
- This ideal  $f(x) = E(Y|X = x)$  is called the regression function.

# The regression function $f(x)$

- Is also defined for vector  $X$ ; e.g.  
 $f(x) = f(x_1, x_2, x_3) = E(Y|X_1 = x_1, X_2 = x_2, X_3 = x_3)$
- Is the **ideal** or **optimal** predictor of  $Y$  with regard to mean-squared prediction error:  $f(x) = E(Y|X = x)$  is the function that minimizes  $E[(Y - g(X))^2|X = x]$  over all functions  $g$  at all points  $X = x$ .
- $\epsilon = Y - f(x)$  is the **irreducible** error – i.e. even if we knew  $f(x)$ , we would still make errors in prediction, since at each  $X = x$  there is typically a distribution of possible  $Y$  values.
- For any estimate  $\hat{f}(x)$  of  $f(x)$ , we have

$$E[(Y - \hat{f}(X))^2|X = x] = \underbrace{[f(x) - \hat{f}(x)]^2}_{\text{Reducible}} + \underbrace{\text{Var}(\epsilon)}_{\text{Irreducible}}$$

## How to estimate $f$

- Typically we have few if any data points with  $X = 4$  exactly.
- So we cannot compute  $E(Y|X = x)$ !
- Relax the definition and let

$$\hat{f}(x) = \text{Ave}(Y|X \in N(x))$$

where  $N(x)$  is some neighborhood of  $x$ .

- Nearest neighbor averaging can be pretty good for small  $p$  – i.e.  $p \leq 4$  and large-ish  $N$ .
- We will discuss smoother versions, such as kernel and spline smoothing later in the course.
- Nearest neighbor methods can be lousy when  $p$  is large. Reason: the **curse of dimensionality**. Nearest neighbors tend to be far away in high dimensions.
  - We need to get a reasonable fraction of the  $N$  values of  $y_i$  to average to bring the variance down – e.g. 10%.
  - A 10% neighborhood in high dimensions need no longer be local, so we lose the spirit of estimating  $E(Y|X = x)$  by local averaging.

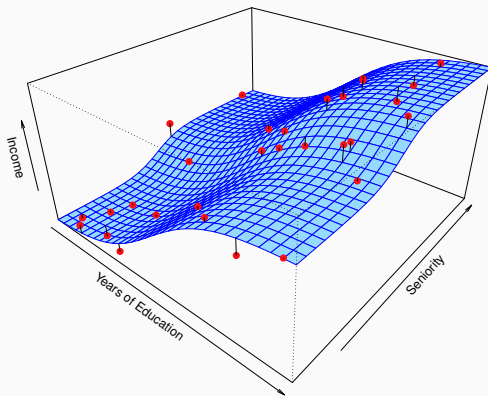
# Parametric and structured models

The **linear** model is an important example of a parametric model:

$$f_L(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p.$$

- A linear model is specified in terms of  $p + 1$  parameters  $\beta_0, \beta_1, \beta_2, \dots, \beta_p$ .
- We estimate the parameters by fitting the model to training data.
- Although it is **almost never correct**, a linear model often serves as a good and interpretable approximation to the unknown true function  $f(X)$ .

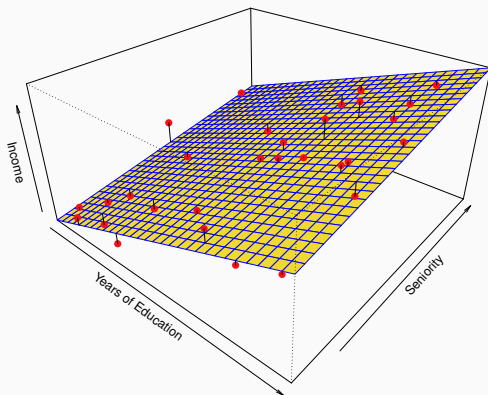
# Simulated example



Red points are simulated values for `income` from the model

$$income = f(education, seniority) + \epsilon$$


$f$  is the blue surface.



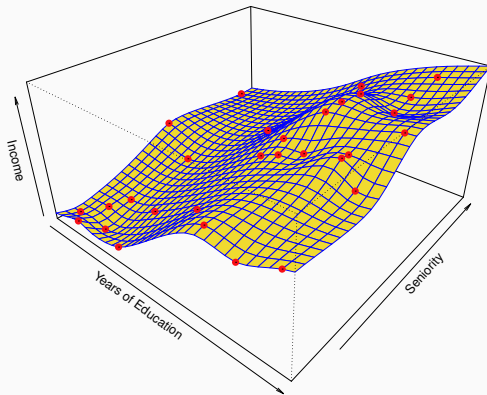
Linear regression model fit to the simulated data.

$$\hat{f}_L(\text{education}, \text{seniority}) = \hat{\beta}_0 + \hat{\beta}_1 \times \text{education} + \hat{\beta}_2 \times \text{seniority}$$





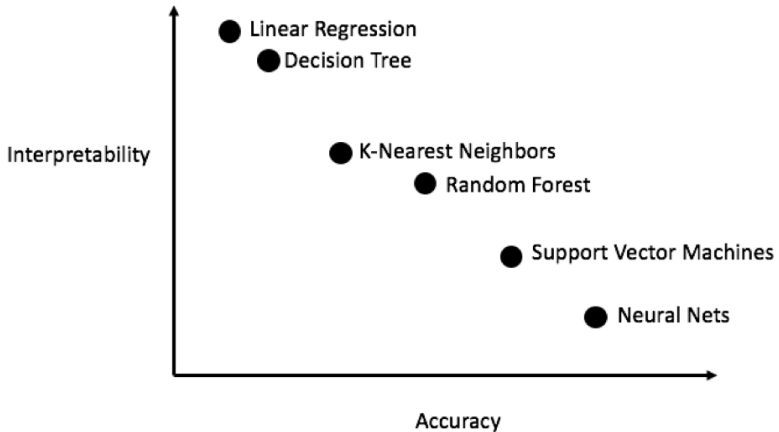
- More flexible regression model  $\hat{f}_S(\text{education}, \text{seniority})$  fit to the simulated data.
- Here we use a technique called a **thin-plate spline** to fit a flexible surface.
- We control the roughness of the fit.



- Even more flexible spline regression model  $\hat{f}_S(\text{education}, \text{seniority})$  fit to the simulated data.
- Here the fitted model makes no errors on the training data!
- Also known as **overfitting**.

## Some trade-offs

- Prediction accuracy versus interpretability.
  - Linear models are easy to interpret; thin-plate splines are not.
- Good fit versus over-fit or under-fit.
  - How do we know when the fit is just right?
- Parsimony versus black-box.
  - We often prefer a simpler model involving fewer variables over a black-box predictor involving them all.



Source: (<https://medium.com/ansaro-blog/interpreting-machine-learning-models-1234d735d6c9>)

# Assessing Model Accuracy

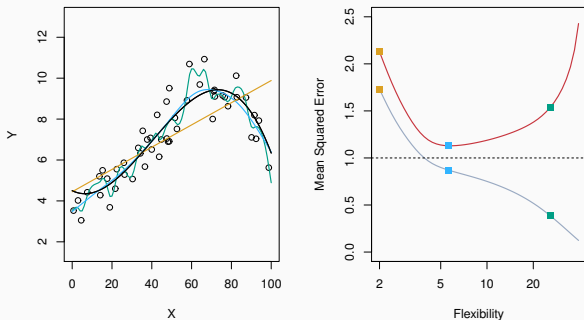
- Suppose we fit a model  $\hat{f}(x)$  to some training data  $Tr = \{x_i, y_i\}_1^N$ , and we wish to see how well it performs.
- We could compute the average squared prediction error over  $Tr$ :

$$MSE_{Tr} = Ave_{i \in Tr} [y_i - \hat{f}(x_i)]^2$$

This may be biased toward more overfit models.

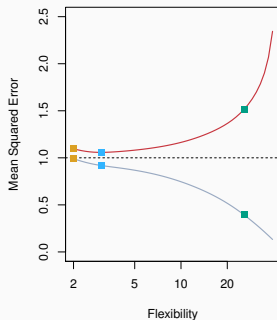
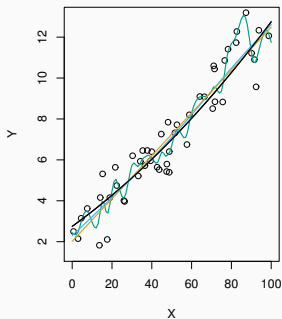
- Instead we should, if possible, compute it using fresh **test** data  $Te = \{x_i, y_i\}_1^M$ :

$$MSE_{Te} = Ave_{i \in Te} [y_i - \hat{f}(x_i)]^2$$

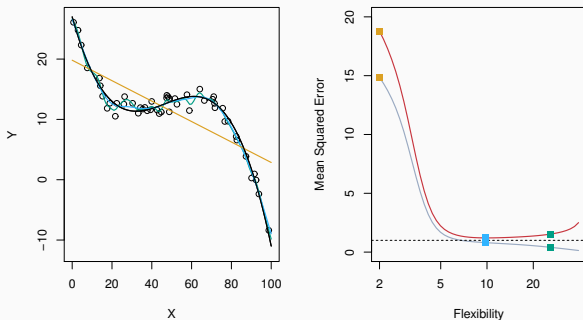


Data simulated from  $f$ , shown in black. Three estimates of  $f$  are shown: the linear regression line (orange curve), and two smoothing splines.

- Black curve is truth.
- Red curve on right is  $MSE_{Te}$ , grey curve is  $MSE_{Tr}$ .
- Orange, blue and green curves/squares correspond to fits of different flexibility.



- The setup as before, using a different true  $f$  that is much closer to linear. In this setting, linear regression provides a very good fit to the data.
- Here the truth is smoother, so the smoother fit and linear model do really well.



- Setup as above, using a different  $f$  that is far from linear.
- In this setting, linear regression provides a very poor fit to the data.
- Here the truth is wiggly and the noise is low, so the more flexible fits do the best.



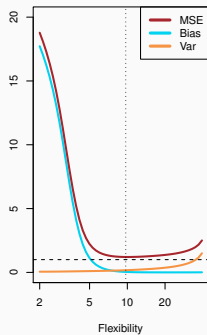
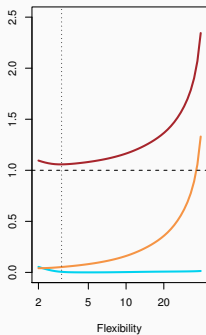
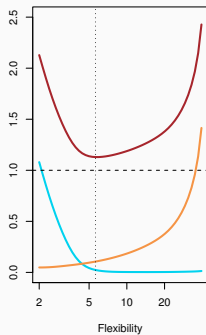
# Bias-Variance Trade-off

- Suppose we have fit a model  $f(x)$  to some training data  $Tr$ , and let  $(x_0, y_0)$  be a test observation drawn from the population.
- If the true model is  $Y = f(X) + \epsilon$  (with  $f(x) = E(Y|X = x)$ ), then

$$E(y_0 - \hat{f}(x_0)) = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon).$$

- The expectation averages over the variability of  $y_0$  as well as the variability in  $Tr$ . Note that  $\text{Bias}(\hat{f}(x_0)) = E[\hat{f}(x_0)] - f(x_0)$ .
- Typically as the **flexibility** of  $\hat{f}$  increases, its variance increases, and its bias decreases.
- So choosing the flexibility based on average test error amounts to a **bias-variance trade-off**.

# Bias-variance trade-off for the three examples



# Classification Problems

Here the response variable  $Y$  is **qualitative** – e.g. email is one of  $\mathcal{C} = (\textit{spam}, \textit{ham})$  ( $\textit{ham} = \textit{godemail}$ ), digitclass is one of  $\mathcal{C} = \{0, 1, \dots, 9\}$ . Our goals are to:

- Build a classifier  $\mathcal{C}(X)$  that assigns a class label from  $\mathcal{C}$  to a future unlabeled observation  $X$ .
- Assess the uncertainty in each classification.
- Understand the roles of the different predictors among  $X = (X_1, X_2, \dots, X_p)$ .

- Is there an ideal  $\mathcal{C}(X)$ ? Suppose the  $K$  elements in  $\mathcal{C}$  are numbered  $1, 2, \dots, K$ . Let

$$p_k(x) = \Pr(Y = k|X = x), k = 1, 2, \dots, K.$$

- These are the **conditional class probabilities** at  $x$ . Then the **Bayes optimal** classifier at  $x$  is

$$\mathcal{C}(x) = j \text{ if } p_j(x) = \max\{p_1(x), p_2(x), \dots, p_K(x)\}$$

- Nearest-neighbor averaging can be used as before.
- Also breaks down as dimension grows.
- However, the impact on  $\hat{\mathcal{C}}(x)$  is less than on  $\hat{p}_k(x)$ ,  $k = 1, \dots, K$ .

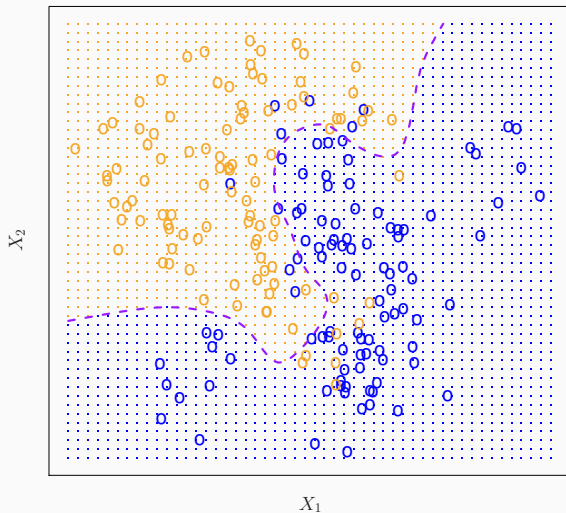
## Classification: some details

- Typically we measure the performance of  $\hat{\mathcal{C}}(x)$  using the misclassification error rate:

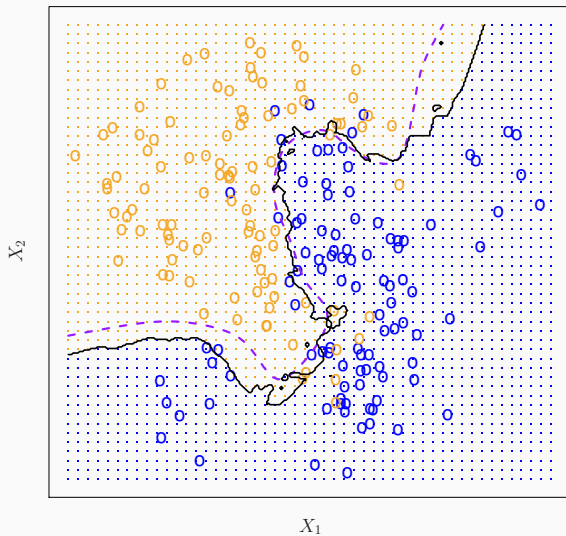
$$Err_{Te} = Ave_{i \in Te} \mathcal{I}[y_i \neq \hat{\mathcal{C}}(x_i)]$$

- The Bayes classifier (using the true  $p_k(x)$ ) has smallest error (in the population).
- Support-vector machines build structured models for  $\mathcal{C}(x)$ .
- We will also build structured models for representing the  $p_k(x)$ . For example, logistic regression, generalized additive models.

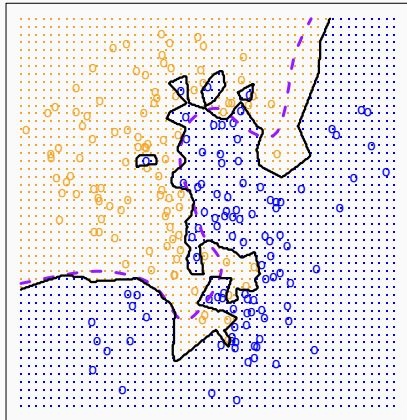
## Example: K-nearest neighbors in two dimensions



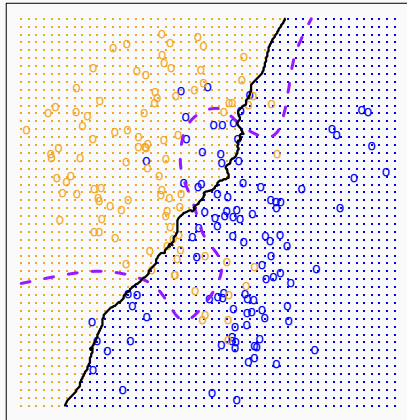
KNN: K=10




KNN:  $K=1$

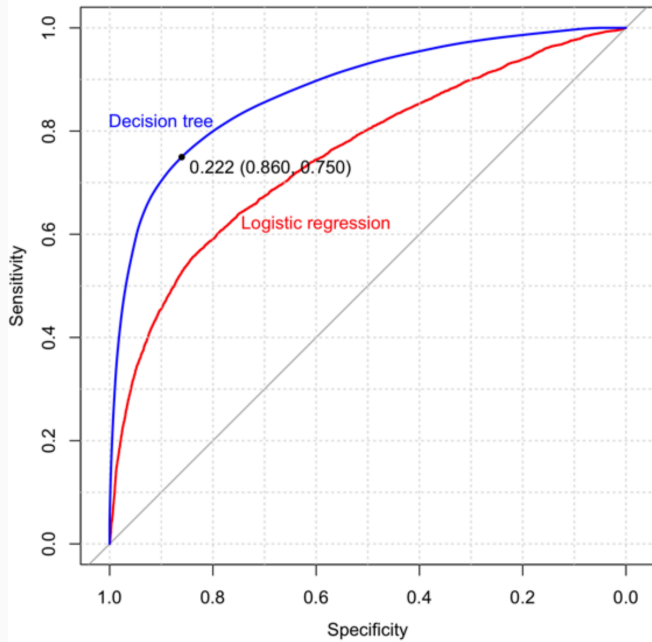


KNN:  $K=100$









# Prediction versus explanation

- Social science: The goal is typically *explanation*
- Data science: The goal is frequently *prediction*, or data exploration
- Many of the same methods are used for both objectives